

## 4. Assignment in “Machine Learning for Natural Language Processing”

Summer Term 2021

### 1 General Questions

1. How are convolutions interpreted in NLP tasks?
2. What is the effect of pooling layers (max/average pooling) in Convolutional Neural Networks? How often are they commonly used?

### 2 Neural Network Hiccups

#### Dying ReLUs

A frequently used activation function for neural networks is the ReLU-function:

$$\text{ReLU}(x) = \begin{cases} x, & x > 0 \\ 0, & \text{else} \end{cases}$$

ReLUs sometimes suffer from the so-called “dying ReLU” problem. In this assignment, you will see what this means and how it can occur.

Assume a neural network with a scalar output, that is, the final layer of the network consists of only one neuron  $o = \sum_{i=1}^n w_i h_i$ , where  $w$  are the weights of the layer and  $h$  is the output of the previous layer. Let’s put a ReLU activation after that layer, to get  $y = \text{ReLU}(o)$ .

We use squared error as the loss function of the network:

$$L = \text{se}(y, t) = \frac{1}{2}(t - y)^2,$$

where  $t$  is the true label for the input.

Let  $w = \begin{pmatrix} 0.2 & -0.3 & 0.5 \end{pmatrix}$ ,  $h = \begin{pmatrix} 0.1 & 0.5 & 10.0 \end{pmatrix}$  and  $t = 0.7$ .

Perform the following steps:

1. Compute the gradient  $\frac{\partial L}{\partial w}$ !
2. Update the weights  $w$  using gradient descent with a learning rate of  $\lambda = 0.1$
3. Repeat steps (1) and (2) with the updated weights, the input  $h = \begin{pmatrix} 0.3 & 0.1 & 1.0 \end{pmatrix}$  and  $t = 0.1$ .
4. Repeat steps (1) and (2) with the updated weights, the input  $h = \begin{pmatrix} 0.5 & 0.25 & 5 \end{pmatrix}$  and  $t = 1$ .

Describe what you find!

## 3 Python

### 3.1 Implementing Neural Networks Part 3 — Dropout

In this assignment, you will implement a neural network “library” yourself, using Python and Numpy (`import numpy as np`). The tool is inspired by PyTorch’s implementation. This week, you will implement Dropout regularisation.

For this, implement a new module (`forward` and `backward` function) that is initialised with a parameter `p`, denoting the probability that a weight is dropped. Do not forget to scale the resulting weights to make up for the missing weights.

```
class Dropout:
    def __init__(self, p=0.5):
        self.p = p

    def forward(self, x: np.array) -> np.array:
        #...

    def backward(self, grad: np.array = np.array([[1]]))
        -> np.array:
        #...
```

Modify the implementation of the `NeuralNetwork` to include dropout with a specified rate (`p: float=0.5` in the constructor) after every hidden layer! Then check that each run of the program will result in different results due to the random cancellation of weights.