

## 8. Assignment in “Machine Learning for Natural Language Processing”

Summer Term 2021

### 1 General Questions

1.
    - What are the problems in Machine Translation with simple RNNs that can be solved by Sequence to Sequence models?
    - What additional problems can be solved with the Attention Mechanism?
- Expected output of different length than input, sometimes words that occur later in the sentence are important for translating a word that occurs earlier.
  - Attention solves the “information bottleneck” produced by encoding the entire input into one vector representation and enables the decoder to focus on the most important parts for each output step.

### 2 Attention

Given a Sequence to Sequence model with an encoder and decoder RNN, that translates text from English to German. Apply the (Luong-)attention mechanism to the hidden states of the encoder RNN.

For the input “I love Language Processing”, the encoder produces the following hidden states  $h_i$ :

$$h_1 = \begin{pmatrix} 1 \\ 0 \\ 1 \end{pmatrix}, h_2 = \begin{pmatrix} 1 \\ 2 \\ 0 \end{pmatrix}, h_3 = \begin{pmatrix} 0 \\ 2 \\ 2 \end{pmatrix}, h_4 = \begin{pmatrix} 0 \\ 2 \\ 1 \end{pmatrix}.$$

The current output of the decoder is “Ich”, the decoder’s hidden state is

$$s_1 = \begin{pmatrix} 2 \\ 1 \\ 0 \end{pmatrix}.$$

Calculate the context vector  $c$ , which is used to predict the next word.

1. First, the dot product of  $s_1$  with all  $h_i$  is calculated.

$$h_1^T s_1 = 2, \quad h_2^T s_1 = 4, \quad h_3^T s_1 = 2, \quad h_4^T s_1 = 2$$

2. These values are then normalized using the softmax function. This results in attention weights of

$$a_1 = 0.09625514, \quad a_2 = 0.71123459, \quad a_3 = 0.09625514, \quad a_4 = 0.09625514$$

3. Finally, the hidden states are weighted using the attention weights and summed up. This results in the context vector

$$\begin{aligned} c &= 0.09625514 \cdot \begin{pmatrix} 1 \\ 0 \\ 1 \end{pmatrix} + 0.71123459 \cdot \begin{pmatrix} 1 \\ 2 \\ 0 \end{pmatrix} + 0.09625514 \cdot \begin{pmatrix} 0 \\ 2 \\ 2 \end{pmatrix} + 0.09625514 \cdot \begin{pmatrix} 0 \\ 2 \\ 1 \end{pmatrix} \\ &= \begin{pmatrix} 0.80748973 \\ 1.80748973 \\ 0.38502054 \end{pmatrix} \end{aligned}$$

### 3 Recap: Backpropagation

Given the network depicted below, perform backpropagation to get the gradients for all weight matrices!

The input is  $x = [1, 0.5, 0.75, 0.25]$ , the target output is  $t = 0$ .

The initial weights are

$$W_1 = \begin{pmatrix} 0.1 & 0.2 & 0.6 & 0.9 \\ 0.7 & 0.5 & 0.3 & 0.6 \\ 0.5 & 0.75 & 0.6 & 0.3 \\ 1. & 0.1 & 0.1 & 0.2 \end{pmatrix}$$

$$W_2 = \begin{pmatrix} 0.6 & 0.6 & 0.6 & 0.2 \\ 1. & 0.3 & 0.1 & 0.1 \\ 0.2 & 0.5 & 0.1 & 0.7 \\ 0.75 & 0.3 & 0.5 & 0.9 \end{pmatrix}$$

$$W_3 = \begin{pmatrix} 0.6 \\ 0.7 \\ 0.2 \\ 0.9 \end{pmatrix}$$

No bias is used in the network. As loss function, we use the squared error,  $L = \frac{1}{2}(t-y)^2$ , where  $t$  is the true label.

Use the vectorised version of backpropagation introduced in the lecture.

First the forward pass:

$$a = w_1^T x = (1.075 \quad 1.0375 \quad 1.225 \quad 1.475)$$

Since all these values are  $> 0$ , the ReLU does "nothing" and

$$b = a.$$




Figure 1: Neural Network with two hidden layers

$$c = w_2^T b = \begin{pmatrix} 3.03375 & 2.01125 & 1.00875 & 2.50375 \end{pmatrix}$$

Again, the ReLU does nothing:

$$\begin{aligned} d &= c \\ e &= w_3^T d = \begin{pmatrix} 5.80325 \end{pmatrix} \\ y &= e \end{aligned}$$

Now for the backward pass, starting at the derivative of the loss by the loss, which is simple:

$$\frac{\partial L}{\partial L} = 1$$

For the loss function, we already know that  $l'(y, t) = y - t$ :

$$\frac{\partial L}{\partial y} = y - t = 5.80325$$

Since the input to the ReLU is  $> 0$ , its local gradient is 1;

$$\frac{\partial L}{\partial e} = \frac{\partial y}{\partial e} \frac{\partial L}{\partial y} = 1 \cdot 5.80325$$

For the next step, we use the formulas from the lecture:

$$\frac{\partial L}{\partial d} = \frac{\partial L}{\partial e} \cdot w_3^T = \begin{pmatrix} 3.48195 & 4.062275 & 1.16065 & 5.22925 \end{pmatrix}$$

$$\frac{\partial L}{\partial w_3} = d^T \frac{\partial L}{\partial e} = \begin{pmatrix} 17.6056 \\ 11.6718 \\ 9.33598 \\ 14.5299 \end{pmatrix}$$

This repeats for the next layers:

$$\frac{\partial L}{\partial c} = \frac{\partial L}{\partial d} \frac{\partial d}{\partial c} = \frac{\partial L}{\partial d}$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial c} \cdot w_2^T = \begin{pmatrix} 6.26751 & 5.33899 & 6.49964 & 9.1111025 \end{pmatrix}$$

$$\frac{\partial L}{\partial w_2} = b^T \frac{\partial L}{\partial c} = \begin{pmatrix} 3.74309625 & 4.36694562 & 1.24769875 & 5.61464438 \\ 3.61252313 & 4.21461031 & 1.20417438 & 5.41878469 \\ 4.26538875 & 4.97628687 & 1.42179625 & 6.39808313 \\ 5.13587625 & 5.99185562 & 1.71195875 & 7.70381438 \end{pmatrix}$$

$$\frac{\partial L}{\partial a} = \frac{\partial L}{\partial b} \cdot 1$$

$$\frac{\partial L}{\partial w_1} = x^T \frac{\partial L}{\partial b} = \begin{pmatrix} 6.26751 & 5.33899 & 6.49964 & 9.1111025 \\ 3.133755 & 2.669495 & 3.24982 & 4.55555125 \\ 4.7006325 & 4.0042425 & 4.87473 & 6.83332688 \\ 1.5668775 & 1.3347475 & 1.62491 & 2.27777563 \end{pmatrix}$$

We do not need to calculate  $\frac{\partial L}{\partial x}$ , since the gradient for the input is not necessary for training.